

Running Wordcount with Hadoop streaming, using Python code

Map/Reduce Module, Running Wordcount with streaming, using Python code

1. Open a Terminal (Right-click on Desktop or click Terminal icon in the top toolbar)
2. Review the following to create the python code

Section 1: wordcount_mapper.py

```
#!/usr/bin/env python
#the above just indicates to use python to intepret this file

# -----
#This mapper code will input a line of text and output <word, 1>
#
# -----

import sys                #a python module with system functions for this OS

# -----
# this 'for loop' will set 'line' to an input line from system
#   standard input file
# -----
for line in sys.stdin:

#-----
#sys.stdin call 'sys' to read a line from standard input,
# note that 'line' is a string object i.e. variable, and it has methods to apply to i
t,
# as in the next line
# -----
```

```

line = line.strip() #strip is a method, ie function, associated
                    # with string variable, it will strip
                    # the carriage return (by default)

keys = line.split() #split line at blanks (by default),
                    # and return a list of keys

for key in keys:    #a for loop through the list of keys
    value = 1

    print('{0}\t{1}'.format(key, value) ) #the {} is replaced by 0th,1st items in
format list

                                #also, note that the Hadoop default is 'tab' separates ke
y from the value

```

Section 2: wordcount_reducer.py

The reducer code has some basic parts, refer to the comments in the code.

```

#!/usr/bin/env python

# -----
#This reducer code will input a line of text and
# output <word, total-count>
# -----

import sys

last_key      = None          #initialize these variables
running_total = 0

# -----
# Loop thru file
# -----

for input_line in sys.stdin:
    input_line = input_line.strip()

    # -----

```

```

# Get Next Word      # -----
    this_key, value = input_line.split("\t", 1) #the Hadoop default is tab separates
key value

                                #the split command returns a list of strings, in this case
into 2 variables

    value = int(value)          #int() will convert a string to integer (this progra
m does no error checking)

# -----
# Key Check part
#   if this current key is same
#       as the last one Consolidate
#   otherwise Emit
# -----
if last_key == this_key:      #check if key has changed ('==' is
#   logical equalilty check
    running_total += value    # add value to running total

else:
    if last_key:              #if this key that was just read in
                                #   is different, and the previous
                                #   (ie last) key is not empy,
                                #   then output
                                #   the previous <key running-count>
        print( "{0}\t{1}".format(last_key, running_total) )
                                # hadoop expects tab(ie '\t')
                                #   separation
    running_total = value      #reset values
    last_key = this_key

if last_key == this_key:
    print( "{0}\t{1}".format(last_key, running_total))

```

NOTE: For Python, please refer to the following code:

Python notes:

```
# 1 indentations are required to indicate blocks of code,  
# 2 all code to be executed as part of some flow control  
#    (e.g. if or for statements) must have the same indentation  
#    (to be safe use 4 space per indentation level, and don't  
#    mix with tabs)  
# 3 flow control conditions have a ':' before  
#    the corresponding block of code  
#
```

Type in the following to open a text editor, and then cut and paste the above lines for wordcount_mapper.py into the text editor, save, and exit. Repeat for wordcount_reducer.py

```
> gedit wordcount_mapper.py
```

```
> gedit wordcount_reducer.py
```

Enter the following to see that the indentations line up as above

```
> more wordcount_mapper.py
```

```
> more wordcount_reducer.py
```

Enter the following to make it executable

```
> chmod +x wordcount_mapper.py
```

```
> chmod +x wordcount_reducer.py
```

Enter the following to see the current directory:

```
> pwd
```

It should be /user/cloudera , or something like that.

3. Create some data:

```
> echo "A long time ago in a galaxy far far away" > /home/cloudera/testfile1
```

```
> echo "Another episode of Star Wars" > /home/cloudera/testfile2
```

4. Create a directory on the HDFS file system (if already exists that's OK):

```
hdfs dfs -mkdir /user/cloudera/input
```

5. Copy the files from local filesystem to the HDFS filesystem:

```
hdfs dfs -put /home/cloudera/testfile1 /user/cloudera/input
```

```
hdfs dfs -put /home/cloudera/testfile2 /user/cloudera/input
```

6. See the files on HDFS

```
hdfs dfs -ls /user/cloudera/input
```

7. Run the Hadoop WordCount example with the input and output specified.

Note that the file paths may differ. The '\ ' just means the command continues on next line.

```
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \  
  -input /user/cloudera/input \  
  -output /user/cloudera/output_new \  
  -mapper /home/cloudera/wordcount_mapper.py \  
  -reducer /home/cloudera/wordcount_reducer.py
```

Hadoop prints out a whole lot of logging or error information. If it runs, see something like the following on the screen scroll by:

```
....
```

```
INFO mapreduce.Job: map 0% reduce 0%
```

```
INFO mapreduce.Job: map 67% reduce 0%
```

```
INFO mapreduce.Job: map 100% reduce 0%
```

```
INFO mapreduce.Job: map 100% reduce 100%
```

```
INFO mapreduce.Job: Job job_1442937183788_0003 completed successfully
```

```
...
```

8. Check the output file to see the results:

```
hdfs dfs -cat /user/cloudera/output_new/part-00000
```

9. View the output directory:

```
hdfs dfs -ls /user/cloudera/output_new
```

Look at the files there and check out the contents, e.g.:

```
hdfs dfs -cat /user/cloudera/output_new/part-00000
```

10. Streaming options:

Try: `hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar --help`

or refer to hadoop.apache.org/docs/r1.2.1/

Change the number of reduce tasks to see its effects. Setting it to 0 will execute no reducer and only produce the map output. (Note the output directory is changed in the snippet below because Hadoop doesn't like to overwrite output)

```
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \  
  -input /user/cloudera/input \  
  -output /user/cloudera/output_new_0 \  
  -mapper /home/cloudera/wordcount_mapper.py \  
  -reducer /home/cloudera/wordcount_reducer.py \  
  -numReduceTasks 0
```

Get the output file from this run, and then upload it:

```
> hdfs dfs -getmerge /user/cloudera/output_new_0/* wordcount_num0_output.txt
```

11. Change the number of reducers to 2.